

Reinforcement Learning

Comprehensive Lecture Notes

Vittorio Cava

A.Y. 2026/27

Contents

I	Lecture 2	2
1	Recap: Multi-Armed Bandits (<i>slides 2–18</i>)	2
1.1	Exploration: the Big Pain of RL (<i>slides 3</i>)	2
1.2	Multi-Armed Bandit (<i>slides 4–5</i>)	2
1.3	Regret (<i>slides 6</i>)	3
1.4	Greedy Algorithm (<i>slides 7</i>)	3
1.5	Explore & Commit Algorithm (<i>slides 8</i>)	4
1.6	Hoeffding Inequality (<i>slides 9–10</i>)	4
1.7	Regret Calculation (<i>slides 11–12</i>)	5
1.8	Regret Decaying (<i>slides 13</i>)	5
1.9	Statistics to Maintain & Confidence (<i>slides 14</i>)	6
1.10	Optimism in the Face of Uncertainty (<i>slides 15</i>)	6
1.11	UCB Algorithm (<i>slides 16</i>)	6
1.12	UCB Algorithm: Regret (<i>slides 17</i>)	7
2	From Multi-Armed to Contextual Bandits (<i>slides 19–24</i>)	8
2.1	From Multi-Armed to Contextual Bandits (<i>slides 19–20</i>)	8
2.2	Contextual Bandits: Interaction (<i>slides 21–23</i>)	8
2.3	Contextual Bandits: Example (<i>slides 24</i>)	9
3	Policy (<i>slides 25</i>)	9
4	Contextual Bandits: Regret (<i>slides 26</i>)	9
5	Explore & Commit Algorithm (<i>slides 27–31</i>)	10
5.1	The algorithm (<i>slides 27</i>)	10
5.2	Unbiased estimates: inverse propensity scoring (<i>slides 27–29</i>)	10
5.3	Fitting a policy (<i>slides 27</i>)	11
5.4	Simulation (<i>slides 30</i>)	11

5.5	Explore & Commit Algorithm: Regret (<i>slides 31</i>)	11
6	ϵ -Greedy (<i>slides 32–35</i>)	12
7	Real World Example (<i>slides 36</i>)	14
8	Bayes Theorem Recall (<i>slides 37</i>)	14
9	Bayesian Bandits (<i>slides 38</i>)	15
10	Gaussian Bayesian Bandits Example (<i>slides 39–40</i>)	15
10.1	Setting (<i>slides 39</i>)	15
10.2	Posterior update (<i>slides 40</i>)	15
11	Gaussian Bayesian Bandits: UCB (<i>slides 41–43</i>)	17
12	Gaussian Bayesian Bandits: Thompson Sampling (<i>slides 44–48</i>)	18
12.1	Probability matching (<i>slides 44–45</i>)	18
12.2	The Gaussian algorithm (<i>slides 46–47</i>)	19
12.3	Simulation (<i>slides 48</i>)	20
13	Bayesian Bandits: Thompson Sampling (<i>slides 49</i>)	20
14	Summary	21
A	Supplementary material (in the companion note, not on the slides)	22
B	Notation map and slide/note discrepancies	23
II	Lecture 2 – Part 2	24
C	Recap (<i>slides 2–11</i>)	24
C.1	From Multi-Armed to Contextual Bandits (<i>slides 3</i>)	24
C.2	Contextual Bandits: Interaction (<i>slides 4</i>)	24
C.3	Contextual Bandits: Regret (<i>slides 5</i>)	24
C.4	Explore & Commit Algorithm (<i>slides 6</i>)	25
C.5	ϵ -Greedy (<i>slides 7</i>)	25
C.6	Bayesian Bandits (<i>slides 8</i>)	26
C.7	Gaussian Bayesian Bandits: UCB (<i>slides 9</i>)	26
C.8	Gaussian Bayesian Bandits: Thompson Sampling (<i>slides 10</i>)	26
D	Sequential Decision Making (<i>slides 12–14</i>)	26
E	Markov Assumption (<i>slides 15–17</i>)	27

F	Markov Decision Process (MDP) (<i>slides 18–19</i>)	28
G	Reward (<i>slides 20</i>)	29
H	Deterministic MDP Example (<i>slides 21</i>)	29
I	Stochastic MDP Example (<i>slides 22</i>)	30
J	Policy (<i>slides 23</i>)	31
K	Trajectory Probability (<i>slides 24</i>)	31
L	State Visitation Probability (<i>slides 25–27</i>)	31
	L.1 Definition (<i>slides 25</i>)	31
	L.2 Examples (<i>slides 26–27</i>)	31
M	Another Example MDP (<i>slides 28</i>)	33
N	Infinite Horizon Discounted Setting (<i>slides 29</i>)	33
O	Value Function (<i>slides 30–31</i>)	34
P	Back to Discount Factor (<i>slides 32</i>)	34
Q	Bellman Equation (<i>slides 33–36</i>)	35
R	Discounted State-Action Distribution (<i>slides 37–38</i>)	36
S	Optimal Policy (<i>slides 39–40</i>)	36
T	Bellman Optimality (<i>slides 41–42</i>)	37
U	Bellman Optimality Example (<i>slides 43–45</i>)	37
V	Bellman Optimality (Theorem 1) (<i>slides 46–49</i>)	38
W	Bellman Optimality (Theorem 2) (<i>slides 50–55</i>)	39
X	Summary	41
A	Appendix: supplementary material (in the companion PDF, not on the slides)	42
	A.1 Discounted sums as expectations under $d_{s_0}^\pi$	42
	A.2 Bellman operators (notation used in the PDF’s proofs)	42

Part I

Lecture 2

How to read these notes

Structure. Sections and subsections follow the *slide deck’s own flow*, and each carries its slide range. The formal content of the companion note (definitions, lemmas, theorems, proofs) is merged at the point where the slides introduce the topic. Material that is in the companion note but has *no* corresponding slide is collected in Appendix A and clearly marked as supplementary.

Boxes. Definitions, theorems, ... are framed in colour; every non-obvious result is followed by a blue **Intuition** box. All figures are original TikZ/pgfplots drawings (no external images). Figures labelled “toy simulation” were produced by small seeded simulations whose data are embedded in the source.

Notation. The slides call the context distribution μ , which clashes with the arm means μ_i ; here it is written ρ (as in the companion note, which also writes z_t , $\mathcal{Z}$ for x_t , $\mathcal{X}$). The full notation map and a list of small slide/note discrepancies are in Appendix B.

1 Recap: Multi-Armed Bandits (*slides 2–18*)

Slides 2–18 recall Lecture 1 (bandits and regret). The formal results used later are Hoeffding’s inequality, the explore-then-commit regret bound and the UCB regret bound; they are cited in the companion note as Theorem 5, Theorem 8 and Corollary 9 of Lecture 1. Short derivations are included here so the notes are self-contained.

1.1 Exploration: the Big Pain of RL (*slides 3*)

An RL agent must *carefully and systematically explore*: remember which states it has visited and try to reach unexplored regions. The core tension is the **exploration–exploitation trade-off**: either make the best decision given the current information (*exploit*) or collect more information (*explore*), i.e. sacrifice something now to gain more later. The slide calls it a chicken-and-egg problem: to know which option is best you must try it, but trying it costs you if it is not the best.

Intuition

Going to your favourite restaurant is the safe choice (exploit); trying the new place round the corner may disappoint tonight but could reveal a new favourite for the next hundred dinners (explore). The longer the horizon, the more an early experiment is worth; with one dinner left, you should simply exploit. Bandits are the simplest model in which this trade-off can be analysed exactly.

1.2 Multi-Armed Bandit (*slides 4–5*)

Definition 1.1 (Multi-armed bandit). A *multi-armed bandit* (MAB) is a simplified MDP with a **single state** and K actions (“arms”) $a_1, \dots, a_K$. Arm i has an unknown reward distribution ν_i with mean $\mu_i = \mathbb{E}_{r \sim \nu_i}[r]$. For $t = 0, \dots, T - 1$ the learner (1) pulls an arm $I_t \in \{1, \dots, K\}$ based on the history and (2) observes an i.i.d. reward $r_t \sim \nu_{I_t}$; the rewards of untried arms are *not* observed.

Intuition

Think of K slot machines. Each pull gives you a payoff *and* information about that one machine only (“bandit feedback”). Because information only comes from the arms you choose, *what you do determines what you learn*; this is why exploration is needed at all. The question left open by the slide is: “what are we trying to optimise exactly?” The answer is the regret.

1.3 Regret (*slides 6*)

Definition 1.2 (Regret). With $\mu^* = \max_{i \in [K]} \mu_i$,

$$\text{Regret}_T = T\mu^* - \sum_{t=0}^{T-1} \mu_{I_t}.$$

It is the total expected reward obtained by pulling the best arm for T rounds minus the total expected reward of the arms actually pulled (the *opportunity loss*).

Intuition

Writing $\Delta_i = \mu^* - \mu_i$ and $N_T(i)$ for the number of pulls of arm i , regret equals $\sum_i \Delta_i N_T(i)$: every pull of a suboptimal arm costs its gap. With means 0.6 and 0.4 and $T = 100$, pulling the worse arm 10 times costs $10 \times 0.2 = 2$ reward units, whereas pulling it 50 times costs 10. A good algorithm therefore wants *few* pulls of bad arms, but it can only know an arm is bad by pulling it.

1.4 Greedy Algorithm (*slides 7*)

Algorithm 1.3 (Greedy). Try each arm once, then commit to the arm with the highest observed reward.

The problem is that a bad arm (low μ_i) may produce a high reward by chance, since $r \sim \nu_i$ is random.

Example 1.4 (Two Bernoulli arms). Arm a_1 pays 1 with probability 0.6 (else 0); arm a_2 pays 1 with probability 0.4. Arm a_1 is clearly better, yet with probability $0.4 \times 0.4 = 16\%$ we observe $(r_1, r_2) = (0, 1)$ and the greedy rule commits to a_2 .

	$r(a_2) = 0$ prob. 0.6	$r(a_2) = 1$ prob. 0.4
$r(a_1) = 1$ prob. 0.6	0.36 picks a_1 ✓	0.24 tie
$r(a_1) = 0$ prob. 0.4	0.24 tie	0.16 picks a_2 ×

Figure 1: Outcomes of the first two pulls in the greedy example (probabilities multiply because the arms are independent).

Intuition

Greedy treats one noisy sample as the truth. If ties are broken uniformly, the probability of locking onto the worse arm is $0.16 + 0.24 = 0.40$ (not just the 16% highlighted on the

slide), and once committed it loses 0.2 per round forever, so the expected regret is about $0.4 \times 0.2 T = 0.08 T$: *linear* in T . Any sensible algorithm must keep sampling long enough to be confident.

1.5 Explore & Commit Algorithm (*slides 8*)

Algorithm 1.5 (Explore-then-commit, MAB). Choose an exploration length N per arm (the slide writes $N = T/K$ as a placeholder; it must satisfy $NK \ll T$, and the optimised value is derived on slide 12).

1. **Explore.** For $k = 1, \dots, K$: pull arm k exactly N times and compute the empirical mean $\hat{\mu}_k = \frac{1}{N} \sum_{i=1}^N r_i$.
2. **Commit.** For $t = NK, \dots, T - 1$: pull $\hat{I} = \arg \max_{i \in [K]} \hat{\mu}_i$.

1.6 Hoeffding Inequality (*slides 9–10*)

To say how wrong $\hat{\mu}$ can be, we need a confidence interval.

Theorem 1.6 (Hoeffding inequality (Lecture 1, Thm. 5)). *Let $X_1, \dots, X_N \in [0, 1]$ be i.i.d. with mean μ . For every $\delta \in (0, 1)$, with probability at least $1 - \delta$,*

$$\left| \frac{1}{N} \sum_{i=1}^N X_i - \mu \right| \leq \sqrt{\frac{\ln(2/\delta)}{2N}} = O\left(\sqrt{\frac{\ln(1/\delta)}{N}}\right).$$

Applying it to all K arms at once with a union bound (replace δ by δ/K) gives, with probability at least $1 - \delta$, $|\hat{\mu}_i - \mu_i| \leq \sqrt{\ln(2K/\delta)/(2N)}$ simultaneously for every arm i .




Figure 2: Left (slide 9): the confidence interval around one empirical mean; the true mean (green) lies inside with probability $1 - \delta$. Right (slide 10): during commitment, all K intervals hold simultaneously, so the width contains $\ln(K/\delta)$. Black dots: $\hat{\mu}_i$ and interval ends.

Intuition

The interval shrinks like $1/\sqrt{N}$: to halve the uncertainty you need four times the data. With $\delta = 0.01$ (“the bound holds with probability 99%”) and $N = 100$ samples, the half-width is $\sqrt{\ln(200)/200} \approx 0.163$; with $N = 1000$ it is ≈ 0.051 . For $K = 10$ arms and $N = 100$, the union bound makes it $\sqrt{\ln(2000)/200} \approx 0.195$: the price of “all arms at once” is only logarithmic in K .

1.7 Regret Calculation (slides 11–12)

Let $\hat{I} = \arg \max_i \hat{\mu}_i$ (empirical best arm) and $I^* = \arg \max_i \mu_i$ (best arm).

Theorem 1.7 (Regret of explore-then-commit, MAB). *Let $L = \ln(2K/\delta)$ (the slides write $\ln(K/\delta)$, absorbing constants into the $O(\cdot)$) and $1 \leq NK \leq T$. With probability at least $1 - \delta$,*

$$\text{Regret}_T = \text{Regret}_{\text{explore}} + \text{Regret}_{\text{exploit}} \leq NK + 2T\sqrt{\frac{L}{N}}.$$

Choosing $N^ = (T\sqrt{L}/K)^{2/3}$ gives $\text{Regret}_T \leq 3T^{2/3}K^{1/3}L^{1/3} = O(T^{2/3}K^{1/3}\ln^{1/3}(K/\delta))$.*

Proof. Let $\varepsilon = \sqrt{L/(2N)}$. By Theorem 1.6 all $|\hat{\mu}_i - \mu_i| \leq \varepsilon$ on an event of probability $\geq 1 - \delta$. On it,

$$\mu_{I^*} \leq \hat{\mu}_{I^*} + \varepsilon \leq \hat{\mu}_{\hat{I}} + \varepsilon \leq \mu_{\hat{I}} + 2\varepsilon,$$

so every committed round costs at most 2ε . Exploration lasts NK rounds with per-round regret ≤ 1 , so $\text{Regret}_{\text{explore}} \leq NK$, and $\text{Regret}_{\text{exploit}} \leq (T - NK)2\varepsilon \leq 2T\sqrt{L/N}$. For the second claim let $f(N) = NK + 2T\sqrt{L}N^{-1/2}$. Then $f'(N) = K - T\sqrt{L}N^{-3/2}$ vanishes at $N^* = (T\sqrt{L}/K)^{2/3}$ and f is convex, so this is the minimiser. From $N^{3/2} = T\sqrt{L}/K$ we get $2T\sqrt{L}N^{*-1/2} = 2KN^*$, hence $f(N^*) = 3KN^* = 3K^{1/3}T^{2/3}L^{1/3}$ (ignoring rounding of N^*). $\square$




Figure 3: The explore/exploit trade-off inside Theorem 1.7 for $T = 10^6$, $K = 5$, $\delta = 0.01$ ($L \approx 6.91$): more exploration costs linearly, but shrinks the commitment error like $1/\sqrt{N}$.

Intuition

Too little exploration makes you commit to the wrong arm; too much wastes rounds on bad arms. The optimum balances the two curves (at N^* the commitment term is exactly twice the exploration term). Numerically, for $T = 10^6$, $K = 5$, $\delta = 0.01$: $N^* \approx 6,513$ pulls per arm, i.e. $\approx 32,600$ exploration rounds (3.3% of T), and a regret bound of $\approx 9.8 \times 10^4$, about 9.8% of the maximum possible. Dividing by T , the *average* regret per round decays like $T^{-1/3}$, i.e. it “approaches 0 as $T \rightarrow \infty$ ” (slide 12), but slowly.

1.8 Regret Decaying (slides 13)

The $T^{2/3}$ rate of explore-then-commit is “kind of slow”. The slide asks for something like $O(\sqrt{T})$ and states that $\sqrt{T}$ is the best possible: it is a *lower bound*, no algorithm can do better in the

worst case (the standard minimax bound is $\Omega(\sqrt{KT})$; it is quoted, not proved, in the lecture). So we look for a new algorithm.




Figure 4: Growth rates compared (slide 13 describes them in words). On a log-log plot a rate T^α is a line of slope α .

Intuition

For $T = 10^6$: $\sqrt{T} = 10^3$ versus $T^{2/3} = 10^4$. A tenfold gap, i.e. per-round regret 0.1% versus 1% (with all constants set to 1).

1.9 Statistics to Maintain & Confidence (*slides 14*)

To compute confidence bounds *while* learning we keep, for each arm i at time t :

$$N_t(i) = \sum_{\tau=0}^{t-1} \mathbf{1}\{I_\tau = i\}, \quad \hat{\mu}_t(i) = \frac{1}{N_t(i)} \sum_{\tau=0}^{t-1} \mathbf{1}\{I_\tau = i\} r_\tau,$$

and, with probability $1 - \delta$ (simultaneously for all arms and all times),

$$|\hat{\mu}_t(i) - \mu_i| \leq \sqrt{\frac{\ln(KT/\delta)}{N_t(i)}}.$$

Intuition

Unlike explore-then-commit, arms now have *different* sample sizes, so each arm gets its *own* interval width. The KT inside the logarithm is a union bound over every (arm, time) pair, because the algorithm will consult these bounds at every round.

1.10 Optimism in the Face of Uncertainty (*slides 15*)

Pick the arm with the highest *upper* confidence bound (the top of its interval).

1.11 UCB Algorithm (*slides 16*)

Algorithm 1.8 (UCB). Pull each arm once (first K rounds). For $t = K, \dots, T - 1$:

1. pick $I_t = \arg \max_{i \in [K]} \left(\hat{\mu}_t(i) + \sqrt{\frac{\ln(KT/\delta)}{N_t(i)}} \right)$;
2. observe the reward and update $N_t, \hat{\mu}_t$.




Figure 5: Optimism (slide 15). Arm 2 has a *lower* empirical mean than arm 3, but its wide interval (few pulls) gives the highest upper bound, so UCB explores it. Green dots: true means.

The square-root *reward bonus* is large for arms not tried many times, which drives exploration.

Intuition

UCB acts as if each arm were as good as it *could plausibly be*. Under-sampled arms get an inflated score and are tried; if the optimism was unjustified, the empirical mean drops and $N_t(i)$ grows, so the bonus shrinks and the arm is abandoned. Nothing is hard-coded about *when* to stop exploring: the schedule emerges from the data. Exploration and exploitation are fused into one rule.

1.12 UCB Algorithm: Regret (*slides 17*)

Theorem 1.9 (Regret of UCB (Lecture 1)). *Let $L' = \ln(2KT/\delta)$. With probability at least $1 - \delta$, UCB satisfies $\text{Regret}_T \leq K + 4\sqrt{KT L'} = \tilde{O}(\sqrt{KT})$.*

Proof. On the event that all intervals hold, write $w_t = \sqrt{L'/N_t(I_t)}$. Then

$$\mu^* \leq \hat{\mu}_t(I^*) + \sqrt{\frac{L'}{N_t(I^*)}} \leq \hat{\mu}_t(I_t) + w_t \leq \mu_{I_t} + 2w_t,$$

so the per-round regret is $\mu^* - \mu_{I_t} \leq 2w_t$ (the first inequality is validity of the interval for I^* , the second is the UCB rule, the third is validity for I_t). Summing over $t \geq K$,

$$\sum_{t=K}^{T-1} \frac{1}{\sqrt{N_t(I_t)}} = \sum_i \sum_{n=1}^{N_T(i)-1} \frac{1}{\sqrt{n}} \leq \sum_i 2\sqrt{N_T(i)} \leq 2\sqrt{K \sum_i N_T(i)} \leq 2\sqrt{KT},$$

by $\sum_{n \leq m} n^{-1/2} \leq 2\sqrt{m}$ and Cauchy-Schwarz. The first K rounds cost at most K . Hence $\text{Regret}_T \leq K + 2\sqrt{L'} \cdot 2\sqrt{KT}$. $\square$

Intuition

Each arm contributes $\sum_{n \leq N} 1/\sqrt{n} \approx 2\sqrt{N}$ to the regret sum, and the total number of pulls is capped at T ; the worst case is spreading pulls evenly, which yields $\sqrt{KT}$. For $K = 5$, $T = 10^6$ the *rate* is $\approx \sqrt{5 \cdot 10^6} \approx 2.2 \times 10^3$ (times $\sqrt{\ln}$ and a constant), versus the $T^{2/3}$ rate of explore-then-commit which is of order 10^4 . (The slide cites an external note for details, https://wensun.github.io/CS4789_data/UCB_note_new.pdf.)

2 From Multi-Armed to Contextual Bandits (*slides 19–24*)

2.1 From Multi-Armed to Contextual Bandits (*slides 19–20*)

A multi-armed bandit is *stateless*: an action leads to a reward. A **contextual bandit** adds a context (state) that is observed before acting and that influences which action is best.




Figure 6: Slides 19–20 redrawn. In a contextual bandit the reward depends on the *state/context* and the *action* (the curved arrow).

2.2 Contextual Bandits: Interaction (*slides 21–23*)

Definition 2.1 (Contextual bandit). A contextual bandit is a distribution ρ over a context set $\mathcal{X}$ together with a mean-reward function $r : \mathcal{X} \times \mathcal{A} \rightarrow [0, 1]$, where $|\mathcal{A}| = K$. For $t = 0, 1, \dots, T - 1$:

1. a context $x_t \sim \rho$ is drawn i.i.d. and *revealed*;
2. the learner picks $a_t \in \mathcal{A}$, possibly at random, using x_t and the history;
3. a reward r_t with $\mathbb{E}[r_t \mid x_t, a_t] = r(x_t, a_t)$ is revealed, and nothing is revealed about the other $K - 1$ actions.

Remark. For simplicity the slides (slide 22) assume *deterministic* rewards $r_t = r(x_t, a_t)$, “as the context is the challenge here”; the companion note stresses that every result below holds with noisy rewards too (the proof of Lemma 5.3 conditions on the noise).




Figure 7: The interaction protocol (slide 23). After the reward, a new round starts: contexts are *not* influenced by past actions, they are sampled i.i.d.

Intuition

The i.i.d. context is what separates a contextual bandit from a full MDP: your action affects your *reward* today but not the *state* you will see tomorrow, so there are no long-term consequences to plan around. It is a bandit with side information. The difficulty that remains is that the best action depends on the context, and you only see the reward of the action you actually took.

2.3 Contextual Bandits: Example (*slides 24*)

A major application is **recommendation systems**: the *context* is user information (age, height, weight, job, ...), the *arms* are items to recommend (news, movies, ...), and each arm has a click-through rate (0/1 reward) that we want to maximise. The question: how do we decide which item to propose next *in a personalised way*?

Example 2.2. Two equally frequent user types and two items, with click-through rates

	news	movies
younger users	0.10	0.30
older users	0.35	0.15

A context-blind bandit sees an average CTR of 0.225 for both items and earns 0.225 per round at best, whereas the personalised rule (movies for younger, news for older) earns $\frac{1}{2}(0.30 + 0.35) = 0.325$.

Intuition

Ignoring context averages away exactly the structure that matters. But learning one separate bandit per context is hopeless when contexts are many or continuous (every user is unique). The remedy, developed next, is to learn a *policy* from a restricted class Π that generalises across contexts.

3 Policy (*slides 25*)

Definition 3.1 (Policy). A policy π is a mapping from (all) states/contexts to actions that determines how the agent selects actions. It can be *deterministic*, $a = \pi(x)$, or *stochastic*, $\pi(a | x)$ (equivalently $a \sim \pi(\cdot | x)$). Given a class Π of policies, the *value* of π is its expected reward, $V(\pi) = \mathbb{E}_{x \sim \rho}[r(x, \pi(x))]$.

Intuition

A policy is a rulebook: “if the user looks like this, show that”. For a finite context set $\mathcal{X}$ with K actions there are $K^{|\mathcal{X}|}$ deterministic policies: with $|\mathcal{X}| = 10$ and $K = 3$ that is $3^{10} = 59,049$, huge, but its *logarithm* is only $10 \ln 3 \approx 10.99$. It is this logarithm, $\ln |\Pi|$, that will govern the regret.

4 Contextual Bandits: Regret (*slides 26*)

The optimal policy in the class is $\pi^* = \arg \max_{\pi \in \Pi} \mathbb{E}_{x \sim \rho} r(x, \pi(x)) = \arg \max_{\pi \in \Pi} V(\pi)$. At every round $a_t = \pi_t(x_t)$ is selected and a reward received.

Definition 4.1 (Contextual regret).

$$\text{Reg}(T) = TV(\pi^*) - \sum_{t=0}^{T-1} V(\pi_t) = T \mathbb{E}_{x \sim \rho}[r(x, \pi^*(x))] - \sum_{t=0}^{T-1} \mathbb{E}_{x \sim \rho}[r(x, \pi_t(x))],$$

where π_t is the policy the learner is using at round t (note: *it changes at every iteration*).

Intuition

The first term is the total expected reward if we always used the best policy; the second is what our sequence of learned policies earns. If $\mathcal{X}$ is a single context, $V(\pi) = r(a)$ and this is exactly the multi-armed regret of Section 1. In the recommendation example the best policy earns 0.325 per round, so a context-blind learner has regret at least $(0.325 - 0.225)T = 0.1T$, linear in T .

5 Explore & Commit Algorithm (*slides 27–31*)

5.1 The algorithm (*slides 27*)

Algorithm 5.1 (Explore-then-commit for contextual bandits). Input: horizon T , policy class Π , exploration budget N .

1. **Explore.** For $t = 0, \dots, N - 1$: observe $x_t \sim \rho$, draw $a_t \sim \text{Unif}(\mathcal{A})$, observe $r_t = r(x_t, a_t)$, and build for x_t an *unbiased estimate* $\hat{\mathbf{r}}_t$ of the reward vector $(r(x_t, a))_{a \in \mathcal{A}}$.
2. **Fit.** Compute the policy $\hat{\pi} \in \arg \max_{\pi \in \Pi} \sum_{t=0}^{N-1} \hat{\mathbf{r}}_t[\pi(x_t)]$.
3. **Commit.** For $t = N, \dots, T - 1$: observe $x_t \sim \rho$ and play $a_t = \hat{\pi}(x_t)$.

The missing ingredient is the unbiased estimate $\hat{\mathbf{r}}_t$: we only observe the reward of *one* action.

5.2 Unbiased estimates: inverse propensity scoring (*slides 27–29*)

An estimator $\hat{\theta}$ is *unbiased* if $\mathbb{E}[\hat{\theta}] = \theta$, i.e. $\text{bias}(\hat{\theta}) = \mathbb{E}[\hat{\theta}] - \theta = 0$ (it is centred on the true value). The goal here: $\mathbb{E}_{a_t \sim p}[\hat{\mathbf{r}}_t[a]] = r(x_t, a)$ for all a .

Definition 5.2 (Inverse propensity scoring (IPS)). Suppose that at round t the action was drawn from a *known* distribution p_t with $p_t(a) > 0$ for all a . The IPS reward vector is

$$\hat{\mathbf{r}}_t[a] = \frac{r_t \mathbf{1}[a_t = a]}{p_t(a)}, \quad a \in \mathcal{A},$$

i.e. the vector with a single non-zero entry $r_t/p_t(a_t)$ at the played action. Under uniform exploration, $p_t(a) = 1/K$ and $\hat{\mathbf{r}}_t[a] = \{Kr_t \text{ if } a = a_t, 0 \text{ otherwise}\}$ (slide 29: “ $r_t/(1/|\mathcal{A}|)$ ”).

Lemma 5.3 (IPS is unbiased). *For every context x_t and every action a , $\mathbb{E}_{a_t \sim p_t}[\hat{\mathbf{r}}_t[a] \mid x_t] = r(x_t, a)$. Consequently $\mathbb{E}[\hat{\mathbf{r}}_t[\pi(x_t)]] = V(\pi)$ for every policy π .*

Proof. Condition on x_t and take the expectation over $a_t \sim p_t$ and the reward noise:

$$\mathbb{E}\left[\frac{r_t \mathbf{1}[a_t = a]}{p_t(a)}\right] = \sum_{a' \in \mathcal{A}} p_t(a') \frac{\mathbb{E}[r_t \mid x_t, a'] \mathbf{1}[a' = a]}{p_t(a)} = p_t(a) \frac{r(x_t, a)}{p_t(a)} = r(x_t, a).$$

Averaging over $x_t \sim \rho$ gives $\mathbb{E}[\hat{\mathbf{r}}_t[\pi(x_t)]] = V(\pi)$. □

Intuition

The division by $p_t(a)$ “removes the asymmetry due to sampling” (slide 28): an action played half as often has its rewards counted twice as heavily, exactly like survey weights that up-weight under-sampled groups. A concrete check with $K = 2$, uniform exploration and $r(x, 2) = 0.8$:

whenever arm 2 is played (half the time) we record $\hat{\mathbf{r}} = (0, 1.6)$, otherwise $(\cdot, 0)$, so the average is $\frac{1}{2} \cdot 1.6 + \frac{1}{2} \cdot 0 = 0.8$. The estimator is unbiased but *noisy*: a single observation is either 0 or Kr_t , far from the true mean (made precise in Lemma A.1, Appendix A). Importantly, $\hat{\mathbf{r}}_t$ is defined for *all* actions even though only one is observed, which is what lets us evaluate policies that would have acted differently.

5.3 Fitting a policy (*slides 27*)

Step 2 of Algorithm 5.1 maximises the *counterfactual value estimate* $\hat{V}(\pi) = \frac{1}{N} \sum_{t < N} \hat{\mathbf{r}}_t[\pi(x_t)]$ over Π . By Lemma 5.3 each summand has mean $V(\pi)$, so $\hat{V}(\pi)$ estimates $V(\pi)$ for *every* $\pi \in \Pi$ from the *same* exploration data.

Intuition

Under uniform exploration, a round only “counts” for policy π if $\pi(x_t)$ happens to equal the randomly played a_t (probability $1/K$), and then it counts with weight Kr_t . In the toy problem below (two contexts, two arms) the four deterministic policies have values $V(a_1, a_1) = 0.525$, $V(a_1, a_2) = 0.675$, $V(a_2, a_1) = 0.425$, $V(a_2, a_2) = 0.575$ (red $\rightarrow$ first entry, blue $\rightarrow$ second), so $\pi^* = (a_1, a_2)$. Because Π here contains *all* context-to-action maps, the maximisation splits into one independent argmax per context; for a restricted class (e.g. linear policies) it does not.

5.4 Simulation (*slides 30*)

Slide 30 shows an animation of the explore phase: a timeline of contexts (red/blue) and, per context, the running estimates of each arm’s reward. Here is an original static version.




Figure 8: Toy simulation of the explore phase (seeded; $N = 80$, uniform exploration, two equally likely contexts, deterministic rewards $r(\text{red}, \cdot) = (0.65, 0.45)$, $r(\text{blue}, \cdot) = (0.40, 0.70)$). Solid: per-context IPS averages; dashed: true values. They are noisy at first but settle on the correct ordering, so $\hat{\pi} = \pi^*$ in this run (about 81% of 2000 seeds at $N = 80$).

5.5 Explore & Commit Algorithm: Regret (*slides 31*)

The slide states that the regret “also depends on the size of the class of policies”. Here is the precise statement.

Theorem 5.4 (Regret of contextual explore-then-commit). *Let Π be finite and run Algorithm 5.1 with $1 \leq N \leq T$. With probability at least $1 - \delta$,*

$$\text{Reg}(T) \leq N + TK \sqrt{\frac{2 \ln(2|\Pi|/\delta)}{N}},$$

and choosing N to minimise the right-hand side gives

$$\text{Reg}(T) \leq 3 \cdot 2^{-1/3} T^{2/3} K^{2/3} \ln^{1/3}\left(\frac{2|\Pi|}{\delta}\right).$$

Proof. Write $\hat{V}(\pi) = \frac{1}{N} \sum_{t \leq N} \hat{\mathbf{r}}_t[\pi(x_t)]$. By Lemma 5.3 the summands are i.i.d. with mean $V(\pi)$ and by Lemma A.1 (Appendix A) they lie in $[0, K]$. Hoeffding's inequality (Theorem 1.6, rescaled by K) gives, for one fixed π and any $\delta' \in (0, 1)$,

$$\mathbb{P}\left(|\hat{V}(\pi) - V(\pi)| > K \sqrt{\ln(2/\delta')/(2N)}\right) \leq \delta'.$$

Take $\delta' = \delta/|\Pi|$ and a union bound over the $|\Pi|$ policies: with probability at least $1 - \delta$, every $\pi \in \Pi$ satisfies $|\hat{V}(\pi) - V(\pi)| \leq \varepsilon$ with $\varepsilon = K \sqrt{\ln(2|\Pi|/\delta)/(2N)}$. On that event, the same three-step chain as in Theorem 1.7 applies to $\hat{\pi}$, which maximises $\hat{V}$:

$$V(\pi^*) \leq \hat{V}(\pi^*) + \varepsilon \leq \hat{V}(\hat{\pi}) + \varepsilon \leq V(\hat{\pi}) + 2\varepsilon.$$

Exploration contributes at most N (each round's regret is at most 1); commitment contributes at most $(T - N) \cdot 2\varepsilon \leq 2T\varepsilon$. Adding them gives the first display.

For the second, set $c = \sqrt{2 \ln(2|\Pi|/\delta)}$, so the bound reads $f(N) = N + TKcN^{-1/2}$. As in the MAB case, f is convex and $f'(N) = 1 - \frac{1}{2}TKcN^{-3/2}$ vanishes at $N^* = (TKc/2)^{2/3}$. Writing $u = (TKc/2)^{1/3}$ gives $f(N^*) = u^2 + 2u^2 = 3(TKc/2)^{2/3}$, and substituting $c^{2/3} = 2^{1/3} \ln^{1/3}(2|\Pi|/\delta)$ gives the claim. $\square$

Intuition

It is the same trade-off as in the MAB: exploring costs N in regret; each extra sample sharpens *all* policy estimates simultaneously, and the union bound over policies only costs a $\ln |\Pi|$ factor inside the square root. Numerically, for $|\Pi| = 10^6$, $K = 3$, $\delta = 0.05$, $T = 10^6$: $\ln(2|\Pi|/\delta) \approx 17.5$, the optimal exploration length is $N^* \approx 4.3 \times 10^4$ rounds (4.3% of T) and the bound is $\approx 1.29 \times 10^5$ (12.9% of T). Making the class a million times richer, from $|\Pi| = 10^6$ to 10^{12} , multiplies the bound only by $(3.15/2.60) \approx 1.21$ (the factor $\ln^{1/3}(2|\Pi|/\delta)$ goes $2.02 \rightarrow 2.60 \rightarrow 3.15$ for $|\Pi| = 10^2, 10^6, 10^{12}$). The price of context is thus *logarithmic in the number of policies* but, because uniform exploration wastes a factor K per observation, it grows with K .

Remark (On the exponent of K). Slide 31 prints the bound as $O(T^{2/3} K^{1/3} \ln^{1/3}(|\Pi|))$, while the proof above (which only uses the *range* $[0, K]$ of the IPS terms) yields $K^{2/3}$. The $K^{1/3}$ form is the known sharper rate: using the *variance* bound of Lemma A.1 with Bernstein's inequality instead of Hoeffding's gives deviations of order $\sqrt{K \ln(|\Pi|/\delta)/N}$ and hence $T^{2/3} K^{1/3} \ln^{1/3}(|\Pi|/\delta)$ plus a lower-order term. This refinement is a standard fact but is *not* proved in the companion note, so the theorem is stated with the exponent that is actually proved.

6 ϵ -Greedy (slides 32–35)

Instead of fixing a threshold for exploring and then committing, we can *interleave* exploration and exploitation.

Algorithm 6.1 (ϵ -greedy for contextual bandits). For $t = 0, 1, \dots, T - 1$:

1. observe $x_t \sim \rho$;
2. draw $a_t \sim p_t = (1 - \epsilon_t) \delta(\pi^t(x_t)) + \epsilon_t \text{Unif}(\mathcal{A})$;
3. observe r_t and form the IPS vector $\hat{\mathbf{r}}_t$ using the propensity $p_t(a_t)$ *actually used*;
4. update the policy: $\pi^{t+1} \in \arg \max_{\pi \in \Pi} \sum_{i \leq t} \hat{\mathbf{r}}_i[\pi(x_i)]$.

Gradually **decay** ϵ_t towards zero (slide 34, “Step 3”).

Here $\delta(\pi^t(x_t))$ is the Dirac delta placing all the mass on the action $\pi^t(x_t)$, so explicitly $p_t(a) = (1 - \epsilon_t) \mathbf{1}[a = \pi^t(x_t)] + \epsilon_t/K$. The two extremes are $\epsilon = 0$ (pure exploitation) and $\epsilon = 1$ (uniform exploration). Since $p_t(a) \geq \epsilon_t/K > 0$, IPS is well defined, and Lemma 5.3 applies at every round (it only needs p_t to be known, and p_t may change with t).




Figure 9: Left (slide 33): the Dirac delta used in p_t . Right: the resulting mixture: every arm gets a floor ϵ/K of exploration probability and the greedy arm additionally gets $1 - \epsilon$.

Intuition

ϵ_t is a dial between “trust what I have learned” and “roll a die”. The floor ϵ_t/K is what guarantees every action keeps being sampled so that IPS stays valid. The dial must be turned down *gradually*: from $r_t \leq 1$ and $p_t(a) \geq \epsilon_t/K$ we get $0 \leq \hat{\mathbf{r}}_t[a] \leq K/\epsilon_t$, so if ϵ_t collapses, the estimates of rarely played actions become extremely noisy (e.g. $\epsilon_t = 0.01$, $K = 3$ allows values up to 300). Note that, unlike explore-then-commit, learning never stops, so the policy keeps adapting. (Neither the slides nor the companion note state a regret bound for ϵ -greedy; slide 35 only shows a simulation.)




Figure 10: Toy simulation of ϵ -greedy (seeded, $T = 90$): early on most rounds are exploratory (open red circles); as ϵ_t decays the learner mostly exploits its current policy, which here becomes $\pi = (\text{red} \rightarrow \text{arm 1, blue} \rightarrow \text{arm 2}) = \pi^*$ (27 of the 90 rounds were exploratory). This ends optimal in about 71% of 1000 seeds at $t = 90$, so individual runs vary; same colours and dashed true values as the earlier figure.

7 Real World Example (*slides 36*)

Slide 36 points to a production use of contextual-bandit algorithms: Microsoft's *Azure Personalizer* (<https://azure.microsoft.com/en-us/products/cognitive-services/personalizer/>), which the slide notes is *going out of service in October 2026* (the slide shows a screenshot of the product page, not reproduced here).

Intuition

Personalising content maps one-to-one onto the model of this lecture: the context is the user/page features, the actions are the candidate items, and the reward is a click or other feedback. Each shown item produces feedback for *that* item only, which is why IPS-style estimates and exploration are needed in practice. Whether a particular product is still available is a separate matter from the algorithms, which remain standard.

8 Bayes Theorem Recall (*slides 37*)

Theorem 8.1 (Bayes' theorem). *For a hypothesis H and data D with $p(D) > 0$,*

$$p(H | D) = \frac{p(H) p(D | H)}{p(D)}.$$




Figure 11: The annotated Bayes formula of slide 37.

Intuition

Posterior $\propto$ prior $\times$ likelihood; $p(D)$ only renormalises. Example: an ad is either *good* (CTR 0.6) or *bad* (CTR 0.4), prior 50/50. After observing one click, $p(\text{good} \mid \text{click}) = \frac{0.5 \cdot 0.6}{0.5 \cdot 0.6 + 0.5 \cdot 0.4} = 0.6$; after two clicks the odds multiply by 0.6/0.4 again, giving odds 2.25, i.e. $p \approx 0.69$. Each observation nudges belief towards the hypotheses that explain it best. This is the engine of Bayesian bandits.

9 Bayesian Bandits (*slides 38*)

So far we made *no assumption* about the reward distributions ν_i ; we only derived bounds that hold for any bounded rewards. In **Bayesian bandits**:

- we exploit *prior* knowledge of the rewards;
- we update a *posterior* distribution of rewards from the history;
- we use the posterior to guide exploration, via **upper confidence bounds** (Bayesian UCB) or **probability matching** (Thompson sampling).

Intuition

A frequentist guarantee protects you against *any* reward distribution but is therefore conservative; a Bayesian method gets a head start when the prior is sensible (e.g. CTRs are mostly small) and the posterior directly tells you “how sure am I about each arm”. The price is that the guarantee is only as good as the modelling assumptions.

10 Gaussian Bayesian Bandits Example (*slides 39–40*)

10.1 Setting (*slides 39*)

Assume each reward distribution is Gaussian, $\nu_i = \mathcal{N}(\mu(i), \sigma^2(i))$. After some pulls the learner holds, for each arm, a belief (a Gaussian) about that arm’s quality: the centre is the best guess $\mu(i)$ and the width reflects remaining uncertainty.

10.2 Posterior update (*slides 40*)

The prior is often $\mathcal{N}(0, 1)$. The posterior after history h_t is proportional to the prior times the product of the likelihoods of all the rewards observed when this arm was selected (slide 40:




Figure 12: Beliefs over the quality of three arms (slide 39 redrawn; dashed lines mark the means $\mu(a_i)$; the double arrows are confidence widths $\pm c\sigma$ around each mean). Arm 1 is very uncertain, arm 3 is well known.

“product over all time-steps t that select action i ”). For a Gaussian prior and Gaussian noise this has a closed form.

Proposition 10.1 (Gaussian conjugacy). *Let the unknown mean of an arm have prior $\mu \sim \mathcal{N}(m_0, s_0^2)$, and let $r_1, \dots, r_n$ be i.i.d. $\mathcal{N}(\mu, \sigma^2)$ with σ^2 known. Then the posterior is $\mathcal{N}(m_n, s_n^2)$ with*

$$\frac{1}{s_n^2} = \frac{1}{s_0^2} + \frac{n}{\sigma^2}, \quad m_n = s_n^2 \left(\frac{m_0}{s_0^2} + \frac{1}{\sigma^2} \sum_{i=1}^n r_i \right).$$

Proof. By Bayes’ rule (Theorem 8.1) the posterior density is proportional to prior times likelihood,

$$p(\mu \mid r_{1:n}) \propto \exp\left(-\frac{(\mu - m_0)^2}{2s_0^2}\right) \prod_{i=1}^n \exp\left(-\frac{(r_i - \mu)^2}{2\sigma^2}\right).$$

The exponent is a quadratic in μ ; collecting terms,

$$-\frac{1}{2} \left[\mu^2 \left(\frac{1}{s_0^2} + \frac{n}{\sigma^2} \right) - 2\mu \left(\frac{m_0}{s_0^2} + \frac{\sum_i r_i}{\sigma^2} \right) \right] + \text{const},$$

which is $-\frac{1}{2s_n^2}(\mu - m_n)^2 + \text{const}$ for exactly the s_n^2 and m_n of the statement. A density proportional to a Gaussian kernel is that Gaussian. $\square$

Corollary 10.2 (One-step update). *If arm i has the posterior $\mathcal{N}(\mu_{t-1}(i), \sigma_{t-1}^2(i))$ and is pulled, yielding r_t , its new posterior is $\mathcal{N}(\mu_t(i), \sigma_t^2(i))$ with*

$$\frac{1}{\sigma_t^2(i)} = \frac{1}{\sigma_{t-1}^2(i)} + \frac{1}{\sigma^2}, \quad \mu_t(i) = \sigma_t^2(i) \left(\frac{\mu_{t-1}(i)}{\sigma_{t-1}^2(i)} + \frac{r_t}{\sigma^2} \right);$$

the posteriors of the other arms are unchanged.

Proof. Posterior updating can be done one observation at a time: the old posterior plays the role of the prior and Proposition 10.1 is applied with $n = 1$. Arms not pulled provide no data about their own mean, so their posteriors stay the same. $\square$

This is the compact update shown on slide 40, $p(\mu_t(i), \sigma_t^2(i) \mid r_t) \propto p(\mu, \sigma^2) \mathcal{N}(r_t \mid \cdot)$: the slide writes the Gaussian likelihood of the new reward with the current posterior parameters as shorthand; the closed form above is what is actually computed.




Figure 13: Toy simulation of Proposition 10.1 ($\sigma = 1$, true mean 0.8, seeded data): the posterior narrows like $1/\sqrt{n}$ (standard deviations 1, 0.71, 0.45, 0.24 for $n = 0, 1, 4, 16$) and its centre drifts from the prior mean 0 towards the truth (posterior means 0.00, 0.25, 0.66 for $n = 1, 4, 16$ in this run).

Intuition

Precisions (inverse variances) *add*: prior precision $1/s_0^2$ plus $1/\sigma^2$ per observation, and the posterior mean is a precision-weighted average of the prior mean and the data. Worked example: prior $\mathcal{N}(0, 1)$, $\sigma^2 = 1$, and $n = 4$ observations summing to 3. Then $s_4^2 = 1/(1 + 4) = 0.2$ and $m_4 = 0.2 \cdot (0 + 3) = 0.6$: the sample mean 0.75 is pulled slightly toward the prior mean 0 (like having one extra “virtual” observation equal to 0). The more data an arm has, the less the prior matters, and a rarely tried arm keeps a wide posterior, which is exactly what the exploration rules below exploit.

11 Gaussian Bayesian Bandits: UCB (*slides 41–43*)

Since we are modelling a distribution, we already have a measure of confidence. What is confidence for a Gaussian? Its *standard deviation*. The slide proposes a UCB rule built from it:

Algorithm 11.1 (Bayesian UCB (Gaussian)). At round t play

$$a_t = \arg \max_{i \in [K]} \left(\mu_t(i) + \frac{c \sigma_t(i)}{\sqrt{N_t(i)}} \right),$$

where $\mu_t(i)$ is the posterior mean, $N_t(i)$ the number of pulls of arm i , $\sigma_t(i)$ the Gaussian’s standard deviation parameter and $c > 0$ a constant that sets how optimistic we are.

The slide’s wording “select the action with highest standard deviation” is shorthand: the score is the mean plus c times the uncertainty of the mean, $\sigma_t(i)/\sqrt{N_t(i)}$, which is the standard error of the arm’s mean (with a weak prior it coincides with the posterior standard deviation from Proposition 10.1).

Intuition

Compare with Algorithm 1.8: the distribution-free width $\sqrt{\ln(KT/\delta)/N}$ is replaced by $c\sigma/\sqrt{N}$, with c playing the role of $\sqrt{\ln(KT/\delta)}$. Under a Gaussian belief, $c = 2$ means that the true mean lies below the score with probability $\Phi(2) \approx 97.7\%$. In the snapshot (means 0.55, 0.10, 0.70; counts 25, 4, 100; $c = 2$) the scores are 0.95, 1.10, 0.90: the well-explored best-looking arm is passed over in favour of the least-explored one, as optimism demands. Slide 43 shows an animation of the rule: as an arm is pulled its bonus shrinks, so the scores converge towards the true means and the choice settles on the best arm.




Figure 14: Snapshot of Algorithm 11.1 with $c = 2$, $\sigma = 1$ (illustrative numbers). Blue: posterior mean; orange: bonus $c\sigma/\sqrt{N}$. Arm 2 has the *lowest* mean but is almost unexplored, so its bonus (1.0) lifts it to the top.

12 Gaussian Bayesian Bandits: Thompson Sampling (*slides 44–48*)

12.1 Probability matching (*slides 44–45*)

Thompson sampling (TS) is a way to do *probability matching*: it selects an action with probability equal (or proportional) to the probability that the action is optimal given what we know, i.e. its *posterior probability of being optimal*. It is optimistic in the face of uncertainty, because uncertain actions have a higher probability of attaining the maximum, and it *uses sampling to avoid the complications of actual probability matching*.

Slide example. You are choosing between two ads. Each step you draw a plausible click-through rate for A and for B from your current beliefs, show the ad with the higher draw, observe the result and update your beliefs; the truly better ad is picked more and more over time.

Algorithm 12.1 (Thompson sampling (posterior sampling)). For $t = 0, 1, \dots, T - 1$:

1. for each arm a , draw $\tilde{\mu}_t(a)$ *independently* from the current posterior of $\mu(a)$;
2. play $A_t \in \arg \max_a \tilde{\mu}_t(a)$;
3. observe r_t and update the posterior of arm A_t only.

In the contextual case the posterior is over the parameters of the *reward model*: step 1 draws one model, and step 2 acts greedily with respect to that draw.

Proposition 12.2 (Thompson sampling is probability matching). *Assume the arms' means are independent under the posterior given the history h_t . Then for every arm a (ties broken by the same rule on both sides)*

$$\mathbb{P}(A_t = a \mid h_t) = \mathbb{P}(a \in \arg \max_{a'} \mu(a') \mid h_t),$$

i.e. the algorithm plays each arm exactly as often as it believes that arm is optimal.

Proof. By independence, drawing $\tilde{\mu}_t(a)$ from each arm's marginal posterior independently is a draw of the whole vector $\tilde{\mu}_t$ from the joint posterior of $(\mu(1), \dots, \mu(K))$ given h_t . The event $\{A_t = a\}$ is the event $\{a \in \arg \max_{a'} \tilde{\mu}_t(a')\}$, whose probability under the joint posterior is by definition the posterior probability that a is an optimal arm. $\square$




Figure 15: Two posteriors with independent arms. A has the higher mean but B is much more uncertain; B’s posterior probability of being the better arm is $\Phi(-0.1/\sqrt{0.1^2 + 0.3^2}) \approx 0.376$, so TS plays B about 37.6% of the time.

Intuition

Why “sample, then argmax” (slide 45) is enough: you never compute $\mathbb{P}(\text{arm } a \text{ is best})$ explicitly; you just *simulate one possible world* consistent with your beliefs and act optimally in it. Over many rounds the frequency with which each arm wins its simulated world is precisely its probability of being optimal. In the figure, for independent Gaussians $\mathbb{P}(\mu_B > \mu_A) = \Phi((m_B - m_A)/\sqrt{s_A^2 + s_B^2}) = \Phi(-0.316) \approx 0.376$ (a Monte Carlo check with 10^6 pairs of draws gives 0.375). The less-known arm keeps getting tried *in proportion to how plausibly it might be best*; once its posterior narrows below the other arm’s, that probability collapses to zero and exploration stops by itself.

12.2 The Gaussian algorithm (*slides 46–47*)

Algorithm 12.3 (Thompson sampling, Gaussian arms). For $t = 0, \dots, T$:

1. for each arm $i = 1, \dots, K$ sample a value $\tilde{\mu}_t(i)$ *independently* from $\mathcal{N}(\mu_{t-1}(i), \sigma_{t-1}^2(i))$ (the slides denote this sampled value by $\hat{\mathbf{r}}_i$; it is *not* the IPS vector of Section 5);
2. pull the arm $I_t = \arg \max_{i \in [K]} \tilde{\mu}_t(i)$;
3. observe the reward r_t ;
4. update the posterior of the pulled arm via Corollary 10.2.

Slide 47 adds two remarks. First, the sampled value is an estimate of the *reward*; in a general MDP it can be replaced by the *Q-function*, i.e. we estimate a *distribution over Q* and sample from it. Second, “this can be done with different distributions as well” (see Section 13).

Intuition

Compared with Bayesian UCB, nothing is added to the mean: uncertainty enters *through the randomness of the sample*. A wide posterior occasionally produces a very high draw and so gets tried; a narrow posterior almost always draws close to its mean. This “randomised optimism” needs no tuning constant c , only a sensible prior.




Figure 16: Toy simulation of Algorithm 12.3 (seeded): two Gaussian arms with true means 0.3 and 0.8 (dashed), noise $\sigma = 1$, prior $\mathcal{N}(0, 1)$, $T = 80$. Arm 1 is tried 13 times and then largely abandoned; arm 2 is played in 67 rounds. Final posteriors: arm 1 $\mathcal{N}(0.22, 0.27^2)$ (still wide, as it was rarely sampled), arm 2 $\mathcal{N}(0.69, 0.12^2)$ (narrow).

12.3 Simulation (slides 48)

13 Bayesian Bandits: Thompson Sampling (slides 49)

Slide 49 repeats the algorithm with *Beta posteriors* for an A/B ad test (click-through rates in $[0, 1]$). Each step: sample $\tilde{\theta}_A, \tilde{\theta}_B$ from the Beta posteriors, pick the argmax, observe a click, and update that arm only.

Proposition 13.1 (Beta–Bernoulli conjugacy). *If $\mu \sim \text{Beta}(\alpha, \beta)$ and $r \sim \text{Bernoulli}(\mu)$, then $\mu \mid r \sim \text{Beta}(\alpha + r, \beta + 1 - r)$. After k successes in n trials from the uniform prior $\text{Beta}(1, 1)$, the posterior is $\text{Beta}(1 + k, 1 + n - k)$, with mean $\frac{k+1}{n+2}$ and variance $\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$ (with $\alpha = 1 + k, \beta = 1 + n - k$).*

Proof. The Beta density is proportional to $\mu^{\alpha-1}(1-\mu)^{\beta-1}$ and the Bernoulli likelihood is $\mu^r(1-\mu)^{1-r}$. Their product is proportional to $\mu^{\alpha+r-1}(1-\mu)^{\beta+(1-r)-1}$, the kernel of $\text{Beta}(\alpha+r, \beta+1-r)$. Iterating n times adds the successes to α and the failures to β ; the moments are those of the Beta distribution. $\square$

Algorithm 13.2 (Thompson sampling, Bernoulli rewards). Keep $(\alpha_i, \beta_i) = (1, 1)$ for each arm i . At each round: draw $\tilde{\theta}_i \sim \text{Beta}(\alpha_i, \beta_i)$ for all i , play $I_t = \arg \max_i \tilde{\theta}_i$, observe the click $c_t \in \{0, 1\}$, and set $\alpha_{I_t} \leftarrow \alpha_{I_t} + c_t, \beta_{I_t} \leftarrow \beta_{I_t} + 1 - c_t$.

Intuition

Beta–Bernoulli is the Bernoulli analogue of Proposition 10.1: “pseudo-counts” of successes and failures. For ad A ($k = 6, n = 10$): mean $7/12 \approx 0.583$, variance ≈ 0.0187 . For ad B ($k = 1, n = 3$): mean $2/5 = 0.4$, variance 0.04 . The probability that a draw for A beats a draw for B is ≈ 0.77 (Monte Carlo, 10^6 draws), so Thompson sampling would show A about 77% of the time and still B about 23% of the time, because B’s few observations leave real doubt. Each click or non-click then moves the corresponding posterior and the shares drift towards the truly better ad. This is the “different distributions” remark of slides 47 and 49: the algorithm only needs a posterior one can *sample* from.




Figure 17: Beta posteriors for an A/B test from the uniform prior (illustrative counts). Ad B has little data, so its posterior is broad.

14 Summary

Method	Setting	How it explores	Assumes	Result in this lecture
Greedy	MAB	try each arm once, commit to the best sample	nothing	linear regret (about $0.08T$ in the Bernoulli example)
Explore-then-commit	MAB	N pulls per arm, then commit	rewards in $[0, 1]$; needs T, K, δ to set N	$O(T^{2/3}K^{1/3}\ln^{1/3}(K/\delta))$ (Thm. 1.7)
UCB	MAB	optimism: $\text{mean} + \sqrt{\ln(KT/\delta)/N_t}$	rewards in $[0, 1]$	$\tilde{O}(\sqrt{KT})$, optimal up to logs (Thm. 1.9)
Explore-then-commit	CB	uniform actions + IPS for N rounds, then $\hat{\pi}$	finite Π	$O(T^{2/3}K^{2/3}\ln^{1/3}(\Pi /\delta))$ as proved (Thm. 5.4); slide: $K^{1/3}$
ϵ -greedy	CB	mixture $(1 - \epsilon_t)\delta_{\pi_t} + \epsilon_t\text{Unif} + \text{IPS}$, $\epsilon_t \downarrow 0$	finite/given Π	no bound stated (simulation only)
Bayesian UCB	MAB (Gaussian)	posterior mean $+ c\sigma_t/\sqrt{N_t}$	Gaussian prior and noise	none stated (rule and snapshot)
Thompson sampling	MAB / CB	sample from posterior, act greedily on the sample	sampleable posterior (Gaussian, Beta, ...)	plays arm a with prob. $\mathbb{P}(a \text{ optimal})$ (Prop. 12.2); no regret bound stated

Table 1: Comparison of the methods covered. “CB” = contextual bandit, “MAB” = multi-armed bandit. “None stated” means neither the slides nor the companion note give a bound.

Key take-aways

- **Exploration is unavoidable, and regret measures its price.** Greedy suffers linear regret; explore-then-commit balances exploration against commitment to get $T^{2/3}$; optimism (UCB) gets the optimal $\sqrt{T}$ rate up to logs, with no schedule to tune.
- **Context changes the object being learned, not the logic.** A contextual bandit asks for a *policy*, with i.i.d. contexts and bandit feedback; regret is measured against the best policy in Π , and the cost of a rich class is only $\ln |\Pi|$ (inside a cube root).




Figure 18: Toy simulation (qualitative): Bernoulli bandit with $K = 3$ arms of means 0.7, 0.5, 0.4, $T = 3000$, regret averaged over 300 seeded runs. At $T = 3000$ the average regrets are about 17 (Thompson), 30 (explore-then-commit), 56 (ϵ -greedy) and 61 (UCB). Constants matter in a short horizon: UCB with the conservative slide bonus $\sqrt{\ln(KT/\delta)/N}$ explores a lot, and explore-then-commit works well here only because $N = 60$ was hand-picked (the theory needs the gaps, or T, K, δ and a bound, to choose it). Asymptotically the ordering of *rates* (linear $\gg T^{2/3} \gg \sqrt{T}$) is what the theorems describe.

- **IPS makes learning from partial feedback possible.** Dividing by the known propensity gives an unbiased estimate of every action’s reward (Lemma 5.3), at the price of variance up to K (uniform exploration) or K/ϵ_t (ϵ -greedy), which is why exploration probabilities must stay positive.
- **Explore-then-commit separates the phases, ϵ -greedy interleaves them.** The first is easy to analyse (Theorem 5.4); the second keeps learning and adapting but needs a decay schedule for ϵ_t .
- **Bayesian bandits turn uncertainty into a distribution you can use.** Conjugate updates (Gaussian, Beta–Bernoulli) are one-line formulas; Bayesian UCB adds a multiple of the posterior width, while Thompson sampling samples from the posterior and thereby plays each arm with its probability of being optimal (Proposition 12.2).

A Supplementary material (in the companion note, not on the slides)

The following result is in the companion note but has no corresponding slide; it supplies the technical ingredient (range of the IPS terms) used in the proof of Theorem 5.4.

Lemma A.1 (What unbiasedness costs). *Under uniform exploration, $0 \leq \hat{r}_t[a] \leq K$ and $\text{Var}(\hat{r}_t[a] \mid x_t) \leq K$.*

Proof. The estimator is $Kr_t \leq K$ on one action and 0 elsewhere, which gives the range. For the variance, $\mathbb{E}[\hat{r}_t[a]^2 \mid x_t] = \frac{1}{K} K^2 \mathbb{E}[r_t^2 \mid x_t, a] \leq K$, since rewards lie in $[0, 1]$, and the variance is at most the second moment. $\square$

Intuition

Unbiasedness is bought with noise: each observation is 0 with probability $1 - 1/K$ and Kr_t with probability $1/K$. With $K = 10$ actions and reward 0.5, a single IPS observation is either 0 or 5 although the target is 0.5; its variance is $10 \cdot 0.25 - 0.25 = 2.25$, versus at most 0.25

for a $[0, 1]$ -valued observation. The *range* bound K is what Hoeffding’s inequality uses in Theorem 5.4; the *variance* bound K is what Bernstein’s inequality would use to sharpen the dependence on K (see the remark after that theorem).

B Notation map and slide/note discrepancies

Concept	Slides	Companion note / these notes
context, context set	$x_t, \mathcal{X}$	$z_t, \mathcal{Z}$ / $x_t, \mathcal{X}$
context distribution	μ	ρ
action at round t	a_t (CB), I_t (MAB)	A_t / a_t, I_t
policy at round t	π^t	π_t
regret	Regret_T	$\text{Reg}(T)$
IPS vector	$\hat{\mathbf{r}}_t[a]$	$\hat{r}_t[a]$ / $\hat{\mathbf{r}}_t[a]$

Small inconsistencies noticed while merging the two sources, and how these notes treat them:

- **Slide 8:** “Set $N = T/K$ ” would use all T rounds for exploration; read it as a placeholder with $NK \ll T$. The optimised N^* is on slide 12 and in Theorem 1.7.
- **Slide 12:** “approaches 0 as T goes to infinity” refers to the *average* regret $\text{Regret}_T/T \sim T^{-1/3}$; the regret itself grows (sub-linearly).
- **Slides 11–12 and 14:** confidence widths are written with $\ln(K/\delta)$ or $\ln(KT/\delta)$, hiding the constant 2 and Hoeffding’s factor $\frac{1}{2}$ inside $O(\cdot)$; the proofs here keep $\ln(2K/\delta)$ and $\ln(2KT/\delta)$.
- **Slide 31 vs. Theorem 5.4:** slide $K^{1/3}$, proved bound $K^{2/3}$; see the remark after the theorem.
- **Slide 42:** “highest standard deviation” should read “highest mean plus c standard errors”, as the displayed formula says.
- **Slides 46–47:** the hat-bold- r symbol there denotes a *sample from the posterior*, unrelated to the IPS vector of slides 27–35.
- **Slide 40:** the recursive update is written with the likelihood evaluated at the current posterior parameters; the closed form is Corollary 10.2.

Part II

Lecture 2 – Part 2

Structure follows the slide deck (slides 1–55); formal statements and proofs are merged in from the companion PDF at the point where the slides introduce the topic. Material that appears only in the PDF is collected in the Appendix.

C Recap (*slides 2–11*)

The lecture opens by recalling the bandit material of the previous lectures (slide 2 opens the recap, slide 11 closes it). The point of the recap is to set up the jump from *one-step* decision making to *sequential* decision making.

C.1 From Multi-Armed to Contextual Bandits (*slides 3*)

A multi-armed bandit (MAB) is stateless: the agent picks an action and receives a reward. A *contextual bandit* (CB) adds back some context (a state): before acting, the agent sees a context x_t , and the reward depends on both the context and the arm.




Figure 19: MAB versus contextual bandit (redrawn from slide 3).

Intuition

Think of a MAB as a single slot machine room with no information about who walks in; a CB is the same room where each customer has a profile x_t and the best machine depends on the profile. The crucial point for what follows: in a CB the context is drawn i.i.d., so *my action does not change the next context*. MDPs remove exactly this restriction.

C.2 Contextual Bandits: Interaction (*slides 4*)

For $t = 0, \dots, T - 1$:

1. a new i.i.d. context $x_t \in \mathcal{X}$ appears;
2. select an action $a_t \in \mathcal{A}$ based on historical information and the context;
3. observe the reward $r(x_t, a_t)$ (context- and arm-dependent).

For simplicity the rewards are assumed deterministic, “as the context is the challenge here”.

C.3 Contextual Bandits: Regret (*slides 5*)

The optimal policy is $\pi^* = \arg \max_{\pi \in \Pi} \mathbb{E}_{x \sim \mu} [r(x, \pi(x))]$. At every iteration $a_t = \pi^t(x_t)$ is chosen, and the regret compares the total expected reward of always using π^* with that of our learned

sequence of policies:

$$\text{Regret}_T = T \mathbb{E}_{x \sim \mu}[r(x, \pi^*(x))] - \sum_{t=0}^{T-1} \mathbb{E}_{x \sim \mu}[r(x, \pi^t(x))].$$

Note that the policies π^t are different at every iteration t .

Intuition

Regret is “how much reward I left on the table by not knowing π^* from day one”. If Regret_T grows sublinearly in T , the *average* regret per round goes to 0: eventually the learner behaves like the optimal policy.

C.4 Explore & Commit Algorithm (*slides 6*)

Algorithm C.1 (Explore & Commit for contextual bandits). 1. For $t = 0, \dots, N-1$ (*explore*): observe $x_t \sim \mu$; sample $a_t \sim \text{Unif}(\mathcal{A})$; observe $r_t = r(x_t, a_t)$; build, for x_t , an unbiased estimate $\hat{r}_t$ of the reward vector, with $\mathbb{E}_{a_t \sim \text{Unif}(\mathcal{A})} \hat{r}_t[a] = r(x_t, a) \forall a$:

$$\hat{r}_t[a] = \begin{cases} 0 & a \neq a_t \\ \frac{r_t}{1/|\mathcal{A}|} & a = a_t. \end{cases}$$

2. Compute the policy $\hat{\pi} = \arg \max_{\pi \in \Pi} \sum_{t=0}^{N-1} \hat{r}_t[\pi(x_t)]$.
3. For $t = N, \dots, T-1$ (*commit*): observe $x_t \sim \mu$ and play $a_t = \hat{\pi}(x_t)$.

With a suitable N the regret is

$$\text{Regret}_T = O\left(T^{2/3} K^{1/3} \ln(|\Pi|)^{1/3}\right), \quad K = |\mathcal{A}|.$$

Intuition

The division by $1/|\mathcal{A}|$ is *importance weighting*: we only see the reward of the arm we actually played, which happens with probability $1/|\mathcal{A}|$, so we inflate it by $|\mathcal{A}|$ to be right on average. A numerical feel for the rate: doubling the horizon multiplies the regret bound by only $2^{2/3} \approx 1.59$, so the per-round regret decays like $T^{-1/3}$.

C.5 ϵ -Greedy (*slides 7*)

Instead of fixing a threshold for exploring and then committing, interleave exploration and exploitation. For $t = 0, \dots, T$: observe $x_t \sim \mu$; draw

$$a_t \sim p_t = (1 - \epsilon) \delta(\pi^t(x_t)) + \epsilon \text{Unif}(\mathcal{A});$$

observe $r_t = r(x_t, a_t)$; build an unbiased estimate $\hat{r}_t$ for x_t (so that $\mathbb{E}_{a \sim p_t} \hat{r}_t[a] = r(x_t, a)$); update the policy $\pi^{t+1} = \arg \max_{\pi \in \Pi} \sum_{\tau \leq t} \hat{r}_\tau[\pi(x_\tau)]$. Here $\epsilon = 0$ means pure exploitation and $\epsilon = 1$ uniform exploration. (The slide only asserts unbiasedness; the usual choice is the importance-weighted estimator $\hat{r}_t[a] = \mathbb{I}[a = a_t] r_t / p_t(a_t)$.)

Intuition

Explore & Commit explores first, then stops learning. ϵ -greedy never stops: with $\epsilon = 0.1$, about one round in ten is a random experiment, and the other nine use the current best guess. It is robust to the case where the early estimates were wrong.

C.6 Bayesian Bandits (*slides 8*)

So far no assumption was made about the reward distributions ν_i ; only bounds on rewards were derived. In Bayesian bandits we (i) exploit *prior* knowledge of rewards, (ii) update a *posterior* distribution from historical information, and (iii) use the posterior to guide exploration, either through upper confidence bounds (Bayesian UCB) or probability matching (Thompson sampling).

C.7 Gaussian Bayesian Bandits: UCB (*slides 9*)

Now that a distribution is modelled, a notion of confidence comes for free: for Gaussians it is the standard deviation. The UCB rule picks

$$a_t = \arg \max_{i \in K} \mu_t(i) + \frac{c \sigma_t(i)}{\sqrt{N_t(i)}},$$

where $N_t(i)$ is the number of pulls of arm i and c a constant.

C.8 Gaussian Bayesian Bandits: Thompson Sampling (*slides 10*)

For $t = 0, \dots, T$: for each arm $i = 1, \dots, K$ sample $\tilde{r}_t(i)$ independently from $\mathcal{N}(\mu_{t-1}(i), \sigma_{t-1}^2(i))$; pull $I_t = \arg \max_i \tilde{r}_t(i)$; observe r_t ; update the posterior $p(\mu_t(i), \sigma_t^2(i) \mid r_t)$. This can be done with other distributions as well. The slide stresses that this is an estimate of the *reward*: in more general MDPs it is replaced by the Q function, i.e. we estimate a distribution over Q .




Figure 20: Thompson sampling with Gaussian posteriors (original illustration of the idea on slide 10; the numbers are illustrative). Dashed lines are the independent draws $\tilde{r}(i)$; the uncertain arm 2 gets the highest draw and is pulled.

Intuition

Thompson sampling explores in proportion to uncertainty without any explicit exploration parameter: an arm with a wide posterior occasionally produces a high draw and gets tried; once its posterior narrows and its mean is low, the draws stop beating the others. This is the “probability matching” idea: pull each arm with the probability that it is the best.

D Sequential Decision Making (*slides 12–14*)

The agent interacts with the environment at discrete time steps, by receiving an observation o_t and a reward r_t from the environment and by taking actions a_t (slide 12).

Such a discrete interaction generates a trajectory, or *history*, at each time step t , which the agent uses to choose its action (slide 13):

$$h_t = (o_0, a_0, r_1, o_1, a_1, \dots, r_t, o_t, a_t) \quad (\text{indexing convention as on the slide}).$$




Figure 21: The agent–world loop (slides 12–14, redrawn).

The *state* is a function of the history, $s_t = f(h_t)$, and it is typically hidden or unknown (slide 14).

Intuition

The history grows without bound; no practical agent can condition on all of it. The state is a *summary* of the history that is supposed to keep what matters for the future. Which summary is a good one is exactly the question the Markov assumption answers.

E Markov Assumption (*slides 15–17*)

Definition E.1 (Markov state). A state representation is Markov if for every t , every history h_t and every action a_t ,

$$p(s_{t+1} \mid s_t, a_t) = p(s_{t+1} \mid h_t, a_t).$$

The future is conditionally independent of the past given the present.

Intuition

If I tell you the current state and my action, nothing in the earlier history can sharpen your prediction of the next state. A chess position is Markov (how we reached it is irrelevant for what can happen next). The position of a moving ball alone is *not*: without its velocity, the past positions carry extra information.



Figure 22: Is this problem Markovian?

Slide 16 asks, over a picture of a robot in a corridor with two pillars, “Is this problem Markovian?” The answer depends on what we put in the state: if the state is only the robot’s position but the dynamics depend on, e.g., its velocity or an unobserved internal variable, past observations help predict the future and the state is *not* Markov; adding the missing variables restores the property.

Slide 17 gives the two extremes:

- a state can *always* be made Markov by taking it equal to the whole history, $s_t = h_t$;
- the best case, used in practice, is that the state is (a sufficient statistic of) the latest

observation, $s_t = o_t$. The state is then said to be *fully observable*.




Figure 23: Two extreme choices of state as a function of the history (original illustration of slide 17).

Intuition

Choosing $s_t = h_t$ is a cheat: it is trivially Markov, but the state space explodes with time. Practical RL wants a small state that is still Markov. Fully observable means the observation already contains everything relevant; this is the setting of the rest of the lecture.

F Markov Decision Process (MDP) (*slides 18–19*)

An MDP is sequential decision making under the Markov assumption: Markovian transition dynamics, full observability, and a (generally) stochastic transition kernel $p(s_{t+1} \mid s_t, a_t)$. Alternative notation: $s_{t+1} \sim p(\cdot \mid s_t, a_t)$ or $s' \sim p(\cdot \mid s, a)$ (slide 19).

Definition F.1 (Markov decision process). A discounted infinite-horizon MDP is a tuple $(\mathcal{S}, \mathcal{A}, p, r, \gamma)$: a set of states, a set of actions, transition dynamics $p(s' \mid s, a)$ satisfying Definition E.1, a reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow [0, R_{\max}]$, and a discount factor $\gamma \in [0, 1)$. Unless stated otherwise $\mathcal{S}$ and $\mathcal{A}$ are finite.

(The slides introduce γ later, in the “Infinite Horizon Discounted Setting” part; the tuple is stated here in full for reference.)




Figure 24: The MDP interaction loop (slides 18–19, redrawn).

Intuition

An MDP is the contextual bandit *with consequences*: the next “context” is the next state s' , and it is drawn from a distribution that depends on the action just taken. Every action therefore has two effects: the reward now, and the situation in which the next decision is made.

G Reward (*slides 20*)

A reward r_t is a scalar signal representing feedback and indicating how well the agent is doing at step t . It is a function of state and action, written $r(s, a)$ (sometimes $r(s', a, s)$). A *cost* is the negative of a reward.

Reward hypothesis (slide 20, flagged as an open question): can all goals be achieved through the maximisation of a numerical reward?

Intuition

The reward is the only way we tell the agent what we want; everything else (what is a good state, a good action) must be *derived* from it. That is why designing rewards is delicate: the agent maximises what we write, not what we meant.

H Deterministic MDP Example (*slides 21*)

The slide shows a small deterministic MDP of a developer's day: states *Wake Up* ($r = 0$), *Netflix* ($r = -2$, with a self-loop), *Code & Debug* ($r = -3$), *Nap* ($r = 1$), *Deploy* ($r = 10$) and *Sleep* ($r = 3$, terminal).




Figure 25: The deterministic MDP of slide 21, redrawn. Rewards are attached to states.

Example H.1 (Worked numbers, not on the slide). Take $\gamma = 0.9$, collect the reward of a state when leaving it, and treat Sleep as terminal (its reward is collected once). Then:

- $V(\text{Sleep}) = 3$, $V(\text{Deploy}) = 10 + 0.9 \cdot 3 = 12.7$, $V(\text{Netflix}) = -2/(1 - 0.9) = -20$ (it can only loop).
- From Code & Debug: $Q(\text{Code}, \text{Deploy}) = -3 + 0.9 \cdot 12.7 = 8.43$, $Q(\text{Code}, \text{Sleep}) = -3 + 0.9 \cdot 3 = -0.3$. The best choice is Deploy: a “bad” step (reward -3) pays off.
- Hence the path $\text{Wake} \rightarrow \text{Code} \rightarrow \text{Deploy} \rightarrow \text{Sleep}$ has discounted return $0 + 0.9 \cdot 8.43 = 7.587$, while $\text{Wake} \rightarrow \text{Netflix}$ forever has return $0 + 0.9 \cdot (-20) = -18$.
- Going $\text{Code} \rightarrow \text{Nap} \rightarrow \text{Wake}$ gives $V(\text{Nap}) = 1 + 0.9 \cdot 7.587 \approx 7.83$ and $Q(\text{Code}, \text{Nap}) \approx 4.05 < 8.43$, so Nap loses.

Intuition

The immediate reward at Code & Debug is negative (-3), yet it is the right place to go through: all the value sits behind it (Deploy = 10). A myopic agent at Wake Up compares -2 (Netflix) with -3 (Code & Debug) and picks Netflix, where it is stuck forever with value -20 ; maximising the *return* sends it through Code & Debug instead. This is the entire reason we define value functions.

I Stochastic MDP Example (*slides 22*)

The slide (taken from Sutton & Barto) shows the *recycling robot*: states **high** and **low** battery; actions **search**, **wait**, **recharge**. The transition table of the slide is:

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	—
low	wait	high	0	—
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	—




Figure 26: Transition graph of the recycling robot (slide 22, redrawn). Black dots are actions; edge labels are (probability, reward).

Example I.1 (Worked numbers, not on the slide). Let $\alpha = 0.7$, $\beta = 0.4$, $r_{\text{search}} = 2$, $r_{\text{wait}} = 1$. At **low**, the expected *immediate* reward of **search** is $0.4 \cdot 2 + 0.6 \cdot (-3) = -1.0$, versus $+1$ for **wait**. Searching with a low battery is risky: with probability 0.6 the robot runs flat, is rescued (reward -3) and ends up at **high**.

Intuition

Stochasticity means the same action from the same state can lead to different next states. Here a single decision (“search with a low battery”) is a gamble, and the right choice depends on the state: at **high** searching is cheap, at **low** it is not. A policy is precisely a rule for making this state-dependent choice.

J Policy (*slides 23*)

A policy π is a mapping from (all) states to actions that determines how the agent selects actions. It can be deterministic, $a = \pi(s)$, or stochastic, $\pi(a | s)$, $p(a | s)$ or $a \sim \pi(\cdot | s)$.

Intuition

A policy is a complete *strategy*, not a plan: it says what to do in *every* state, including the ones I might never reach. In the recycling robot there are $2 \cdot 3 = 6$ deterministic policies (two choices at **high**, three at **low**); in general there are $|\mathcal{A}|^{|\mathcal{S}|}$, which is why we will need structure (Bellman equations) rather than brute-force enumeration.

K Trajectory Probability (*slides 24*)

What is the probability of seeing the trajectory $(s_0, a_0, s_1, a_1, \dots, s_t, a_t)$ at time t when following π from s_0 ? Using the Markov property and the fact that π depends only on the current state:

$$\mathbb{P}^\pi(s_0, a_0, \dots, s_t, a_t) = \pi(a_0 | s_0) p(s_1 | s_0, a_0) \pi(a_1 | s_1) \cdots p(s_t | s_{t-1}, a_{t-1}) \pi(a_t | s_t).$$




Figure 27: Graphical model of a trajectory (slide 24, redrawn).

Intuition

Each factor is one coin flip: the agent's flip (π) followed by the world's flip (p). Numbers: if $\pi(a_0 | s_0) = 0.5$, $p(s_1 | s_0, a_0) = 0.9$ and $\pi(a_1 | s_1) = 1$, then the probability of the length-2 trajectory (s_0, a_0, s_1, a_1) is $0.5 \cdot 0.9 \cdot 1 = 0.45$.

L State Visitation Probability (*slides 25–27*)

L.1 Definition (*slides 25*)

What is the probability of visiting (s, a) at time t when following π from s_0 ? Marginalise the trajectory probability over everything except the last pair:

$$\mathbb{P}_t^\pi(s, a; s_0) = \sum_{a_0, s_1, a_1, \dots, s_{t-1}, a_{t-1}} \mathbb{P}^\pi(s_0, a_0, \dots, s_t = s, a_t = a).$$

Intuition

It answers “if I run the policy many times, in what fraction of the runs am I in (s, a) at time t ?” Plotting it for every state gives a heat map that moves and spreads with t .

L.2 Examples (*slides 26–27*)

Slide 26 shows, for a policy that goes to the goal in a gridworld with 10% slip, the visitation probability estimated over 2000 simulations, together with the probability mass in the (absorbing)





Figure 28: Schematic gridworld (left) and shape of the probability mass at the absorbing goal (right) for a goal-directed policy versus a random one. The curve is an original schematic read off slides 26–27, *not* the data of the slide.

goal as a function of t . Slide 27 shows the same for a *random* policy, on a log scale: the mass at the goal stays essentially 0.

For a policy that goes to the goal, 10% slip,
computed over 2000 simulations




Figure 29: Heat map of $d_t^pi(s)$ for the goal-directed policy (10% slip, 2000 simulations) with the “mass at the goal” curve.




Figure 30: Same for the random policy, probability on log scale.

Intuition

Under the goal-directed policy nothing reaches the goal for the first ~ 15 steps (that is the minimum distance), then the mass climbs towards 1; the missing mass is runs that slipped and are still on the way. The random policy diffuses over the grid so thinly that its goal mass is almost 0 at every time: the visitation distribution is a fingerprint of the policy's behaviour.

M Another Example MDP (*slides 28*)

A robotic arm that must move a ball (slide 28):

- **state:** robot configuration (joint states) and ball position;
- **action:** torque on arm and finger joints;
- **transition:** stochastic, physics plus noise;
- **policy:** mapping from robot state and ball position to torque;
- **cost:** magnitude of the torque and distance to the goal.




Figure 31: Photo of a robotic arm manipulating a ball

Intuition

Same five ingredients, a very different world: continuous states and actions, physics instead of a table of probabilities, and a cost made of two competing terms (effort and distance to goal). The abstract tuple $(\mathcal{S}, \mathcal{A}, p, r)$ covers both the toy table above and a robot.

N Infinite Horizon Discounted Setting (*slides 29*)

So far the MDP is $(\mathcal{S}, \mathcal{A}, p, r)$. We now add a discount factor γ to reason about the long-term effects of a policy, with γ in $[0, 1]$ on the slide: $\gamma = 0$ means “I only care about immediate rewards”, while $\gamma = 1$ means immediate and future rewards are equally important. The slide leaves “How so?” as a question; it is answered in “Back to Discount Factor” below, where we see that $\gamma = 1$ is a bad idea for infinite tasks and the formal definition (Definition F.1) takes $\gamma < 1$.

Intuition

γ is a patience dial. At $\gamma = 0.5$ a reward two steps away is worth 0.25 of the same reward now; at $\gamma = 0.99$ it is worth 0.98.

O Value Function (*slides 30–31*)

We estimate the *goodness* of states and actions by their value, which is also a measure to compare policies. The state-value function of a policy π is the expected discounted return from s (slide 30):

$$\begin{aligned} V^\pi(s_t) &= \mathbb{E}\left[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \mid s_t\right] \\ &= \mathbb{E}\left[\sum_{h=0}^{\infty} \gamma^h r_{t+h} \mid s_t = s, a_h = \pi(s_h), s_{h+1} \sim p(\cdot \mid s_h, a_h)\right]. \end{aligned}$$

Slide 31 adds the action-value function, where the first action is fixed and then π is followed:

$$\begin{aligned} Q^\pi(s_t, a_t) &= \mathbb{E}\left[\sum_{h=0}^{\infty} \gamma^h r_{t+h} \mid (s_t, a_t) = (s, a), \right. \\ &\quad \left. a_h = \pi(s_h), s_{h+1} \sim p(\cdot \mid s_h, a_h)\right]. \end{aligned}$$

Intuition

$V^\pi(s)$: “how good is it to *be* in s if I then follow π ?” $Q^\pi(s, a)$: “how good is it to take a *now* in s and then follow π ?” The difference between the two is one free first decision. Comparing $Q^\pi(s, a)$ across a is how an agent decides which action is better than the policy’s current choice.

P Back to Discount Factor (*slides 32*)

Setting $\gamma = 1$ for infinite tasks is a bad idea. The quantity $\sum_{h=0}^{\infty} \gamma^h$ is a geometric series, equal to $1/(1 - \gamma)$ for $\gamma \in [0, 1)$, so the value of γ approximately determines how many steps ahead we are looking: e.g. $\gamma = 0.99$ corresponds to ~ 100 time steps.

Lemma P.1 (Values are finite, and how finite). *For every policy and every state, $0 \leq V^\pi(s) \leq \frac{R_{\max}}{1 - \gamma}$.*

Proof. Rewards lie in $[0, R_{\max}]$, so $0 \leq \sum_h \gamma^h r_h \leq R_{\max} \sum_{h \geq 0} \gamma^h = R_{\max}/(1 - \gamma)$, the geometric series converging because $\gamma < 1$. Expectations preserve the bounds. $\square$

Intuition

With $\gamma = 0.99$ the weight γ^h falls to $\gamma^{100} \approx 0.37 \approx 1/e$ after about 100 steps; with $\gamma = 0.9$ the same happens after about 10 steps. If rewards are at most $R_{\max} = 1$, the value can never exceed $1/(1 - 0.99) = 100$; at $\gamma = 1$ the bound is infinite and the sum of rewards of an endless task need not exist at all. The discount is what makes “the value of a state” a well-defined finite number.




Figure 32: Left: discount weights γ^h . Right: the effective horizon $1/(1 - \gamma)$ diverges as $\gamma \rightarrow 1$ (original figures for slide 32).

Q Bellman Equation (*slides 33–36*)

The value of a state is expanded in terms of the current reward and the value of the next states according to the policy (slide 33): for a deterministic π ,

$$V^\pi(s_t) = \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \mid s_t] = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s_t, \pi(s_t))} [V^\pi(s')],$$

and likewise for Q (slides 34–36):

$$Q^\pi(s_t, a_t) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s_t, a_t)} [V^\pi(s')].$$

In the first equation r is a function of s and $\pi(s)$; in the second, of s and a . As a result $V^\pi(s) = Q^\pi(s, \pi(s))$ (slide 36).

Theorem Q.1 (Bellman expectation equations). *For every policy π and all s, a :*

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_{a \sim \pi(\cdot \mid s)} [r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s, a)} [V^\pi(s')]], \\ Q^\pi(s, a) &= r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s, a)} [V^\pi(s')], \\ V^\pi(s) &= \mathbb{E}_{a \sim \pi(\cdot \mid s)} [Q^\pi(s, a)]. \end{aligned}$$

For a deterministic policy the first reads $V^\pi(s) = r(s, \pi(s)) + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s, \pi(s))} [V^\pi(s')]$ and the last $V^\pi(s) = Q^\pi(s, \pi(s))$.

Proof. Split the return after the first step using $G_0 = r(s_0, a_0) + \gamma G_1$: $V^\pi(s) = \mathbb{E}^\pi[G_0 \mid s_0 = s] = \mathbb{E}^\pi[r(s_0, a_0) \mid s_0 = s] + \gamma \mathbb{E}^\pi[G_1 \mid s_0 = s]$. Condition the second term on (a_0, s_1) and use the tower rule. By Definition E.1 and the fact that π depends only on the current state, the conditional distribution of the trajectory from s_1 onwards given (s_0, a_0, s_1) is the same as that of a trajectory started at s_1 ; hence $\mathbb{E}^\pi[G_1 \mid s_0 = s, a_0 = a, s_1 = s'] = V^\pi(s')$. Averaging over $s' \sim p(\cdot \mid s, a)$ and then over $a \sim \pi(\cdot \mid s)$ gives the first identity; the argument for Q^π is the same with a_0 fixed. $\square$

Intuition

The equation says: *value now = reward now + discounted value of where I land*. It turns an infinite sum into a one-step consistency condition. Example (deterministic): if $r(s, \pi(s)) = 1$, $\gamma = 0.9$ and the next state always has value $V^\pi(s') = 10$, then $V^\pi(s) = 1 + 0.9 \cdot 10 = 10$. The two lines of the theorem differ only in what is averaged: V averages over my action choice (policy), Q has the first action already fixed and averages only over the environment's randomness.




Figure 33: Backup diagrams for the Bellman expectation equations (original figure for slides 33–36).

R Discounted State-Action Distribution (*slides 37–38*)

Definition R.1 (Discounted state-action distribution).

$$d_{s_0}^\pi(s, a) = (1 - \gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}_h^\pi(s, a; s_0).$$

For an initial distribution μ write $d_\mu^\pi(s, a) = \mathbb{E}_{s_0 \sim \mu}[d_{s_0}^\pi(s, a)]$.

Lemma R.2 (It is a distribution, and that is the point of the $(1 - \gamma)$). $d_{s_0}^\pi$ is a probability distribution over $\mathcal{S} \times \mathcal{A}$.

Proof. Non-negativity is clear. Summing over (s, a) and exchanging the (absolutely convergent) sums, $\sum_{s,a} d_{s_0}^\pi(s, a) = (1 - \gamma) \sum_h \gamma^h \sum_{s,a} \mathbb{P}_h^\pi(s, a; s_0) = (1 - \gamma) \sum_h \gamma^h = 1$, since $\mathbb{P}_h^\pi(\cdot; s_0)$ is a distribution for each h . $\square$

This matches the slide’s remark: “this gives us a probability distribution (remember $\sum_h \gamma^h = 1/(1 - \gamma)$)”. The companion identity expressing discounted sums as expectations under $d_{s_0}^\pi$ is not on the slides and is given in Appendix A.1.

Intuition

Take the visitation probabilities of the previous section, weight time step h by γ^h and sum them: this is “how much total (discounted) time I spend in (s, a) ”. Without the factor $(1 - \gamma)$ the total mass would be $1/(1 - \gamma)$ (e.g. 100 for $\gamma = 0.99$); multiplying by $(1 - \gamma)$ rescales it to 1 so it can be used as a genuine probability distribution. Early time steps get most of the weight, which is the same patience dial as before.

S Optimal Policy (*slides 39–40*)

Definition S.1 (Optimal value and optimal policy). $V^*(s) = \sup_\pi V^\pi(s)$ and $Q^*(s, a) = \sup_\pi Q^\pi(s, a)$. A policy π^* is optimal if $V^{\pi^*}(s) = V^*(s)$ for every s .

Slides 39–40 state that for infinite-horizon MDPs there *always* exists a deterministic policy π^* such that

$$V^{\pi^*}(s) \geq V^\pi(s) \quad \forall s, \pi,$$

i.e. π^* dominates all other policies in each state; it returns optimal actions a^* , and $V^{\pi^*} = V^*$, $Q^{\pi^*} = Q^*$. The existence statement is proved constructively by Theorem V.1 below.

Intuition

It is far from obvious that one policy can be best *in all states simultaneously*; in many optimisation problems the best choice for one objective hurts another. Here it is true, and it is also true that we never need randomness: a table of “best action in each state” suffices. This is what makes the search for optimal behaviour tractable.

T Bellman Optimality (*slides 41–42*)

Theorem T.1 (Bellman optimality equation).

$$V^*(s) = \max_a \left[r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s, a)} V^*(s') \right],$$

$$Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s, a)} \left[\max_{a'} Q^*(s', a') \right].$$

On slide 42 the bracket $r(s, a) + \gamma \mathbb{E}_{s'} V^*(s')$ is annotated as $Q^*(s, a)$, so that $V^*(s) = \max_a Q^*(s, a)$.

Intuition

The Bellman *expectation* equation averages over the policy’s action; the Bellman *optimality* equation replaces that average by a max: the optimal agent simply takes the best action. All the difficulty of looking arbitrarily far into the future is hidden in $V^*(s')$: if I knew the optimal value of every neighbour, choosing the best action would be a one-step problem.

U Bellman Optimality Example (*slides 43–45*)

Assume we know V^* at s' and s'' , the two successors of s reached by actions a' and a'' (slide 43). Try a' , get $r(s, a')$ and compute $Q^*(s, a') = r(s, a') + \gamma V^*(s')$; try a'' and compute $Q^*(s, a'') = r(s, a'') + \gamma V^*(s'')$ (slide 44). Then (slide 45)

$$V^*(s) = \max_{a', a''} \{Q^*(s, a'), Q^*(s, a'')\}.$$




Figure 34: The one-step lookahead of slides 43–45 with illustrative numbers ($\gamma = 0.9$; the numbers are mine, not the slide’s).

Intuition

With the numbers above the action with the *larger immediate reward* (a'' , reward 3) is *not* the best: a' leads to a state worth 5, and $1 + 4.5 > 3 + 1.8$. This is the one-step trick that

everything else in the course rests on: evaluate each action as “reward + discounted value of where it leads” and take the best.

V Bellman Optimality (Theorem 1) (*slides 46–49*)

Theorem V.1 (Theorem 1 — greedy with respect to Q^* is optimal). *Let $\hat{\pi}(s) \in \arg \max_a Q^*(s, a)$. Then $V^{\hat{\pi}} = V^*$, and therefore $\pi^* = \hat{\pi}$ is an optimal policy.*

Proof as shown on the slides (slides 46–48)

Starting from $V^*(s) = \max_a [r(s, a) + \gamma \mathbb{E}_{s'} V^*(s')]$ and unrolling the equation along $\hat{\pi}$:

$$\begin{aligned} V^*(s) &= r(s, \pi^*(s)) + \gamma \mathbb{E}_{s'} V^*(s') \\ &\leq \max_a [r(s, a) + \gamma \mathbb{E}_{s'} V^*(s')] = r(s, \hat{\pi}(s)) + \gamma \mathbb{E}_{s'} V^*(s') \\ &\leq r(s, \hat{\pi}(s)) + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \hat{\pi}(s))} [r(s', \hat{\pi}(s')) + \gamma \mathbb{E}_{s''} V^*(s'')] \\ &\leq \dots \leq \mathbb{E} [r(s, \hat{\pi}(s)) + \gamma r(s', \hat{\pi}(s')) + \dots] = V^{\hat{\pi}}(s), \end{aligned}$$

so $V^{\hat{\pi}} \geq V^*$ (the leftover term $\gamma^k \mathbb{E} V^*(s_k) \rightarrow 0$ since V^* is bounded, Lemma P.1); and $V^* \geq V^{\hat{\pi}}$ holds by definition of V^* as a supremum. Hence $V^{\hat{\pi}} = V^*$ and, by slide 49, $\pi^* = \arg \max_a Q^*(s, a)$ is an optimal policy.

Proof in the companion text (fixed-point argument)

Proof. By Theorem T.1, $V^* = \mathcal{B}V^*$ and $Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s'} [V^*(s')]$ (the operators $\mathcal{B}, \mathcal{B}^{\hat{\pi}}$ are recalled in Appendix A.2). Since $\hat{\pi}(s)$ maximises $Q^*(s, \cdot)$,

$$(\mathcal{B}V^*)(s) = \max_a Q^*(s, a) = Q^*(s, \hat{\pi}(s)) = r(s, \hat{\pi}(s)) + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \hat{\pi}(s))} [V^*(s')] = (\mathcal{B}^{\hat{\pi}}V^*)(s),$$

so V^* is a fixed point of $\mathcal{B}^{\hat{\pi}}$. By Theorem Q.1 so is $V^{\hat{\pi}}$, and both are bounded by Lemma P.1. A bounded fixed point of $\mathcal{B}^{\hat{\pi}}$ is unique. If u and v are two, then for every s

$$|u(s) - v(s)| = \gamma |\mathbb{E}_{s'} [u(s') - v(s')]| \leq \gamma \|u - v\|_{\infty},$$

because the rewards cancel, so $\|u - v\|_{\infty} \leq \gamma \|u - v\|_{\infty}$. The norm is finite and $\gamma < 1$, hence $u = v$. Therefore $V^{\hat{\pi}} = V^*$, and $\hat{\pi}$ attains the supremum of Definition S.1 at every state. $\square$

Intuition

Theorem 1 is the license to *act*: once someone hands me Q^* , I do not need to plan any more; the single-step rule “pick the action with the largest $Q^*(s, \cdot)$ ” is already globally optimal. The slide proof says it as a sandwich: following $\hat{\pi}$ can’t be better than optimal ($V^{\hat{\pi}} \leq V^*$), and by unrolling the optimality equation it can’t be worse either ($V^{\hat{\pi}} \geq V^*$). The book proof says the same thing differently: V^* and $V^{\hat{\pi}}$ solve the same one-policy equation, and that equation has only one bounded solution because each application shrinks distances by $\gamma < 1$.

W Bellman Optimality (Theorem 2) (slides 50–55)

Theorem W.1 (Theorem 2 — the optimality equation has one solution). *If $V : \mathcal{S} \rightarrow \mathbb{R}$ satisfies $V(s) = \max_a [r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s,a)} V(s')] for every s , then $V = V^*$.$*

The slides (51–53) ask us to check that $\|V - V^*\|_\infty = 0$ and then unroll the difference.

Proof. Both V and V^* satisfy the same equation, so for every s ,

$$\begin{aligned} |V(s) - V^*(s)| &= \left| \max_a [r(s, a) + \gamma \mathbb{E}_{s'} V(s')] - \max_a [r(s, a) + \gamma \mathbb{E}_{s'} V^*(s')] \right| \\ &\leq \max_a \left| [r(s, a) + \gamma \mathbb{E}_{s'} V(s')] - [r(s, a) + \gamma \mathbb{E}_{s'} V^*(s')] \right| \\ &= \max_a \gamma \left| \mathbb{E}_{s' \sim p(\cdot|s,a)} [V(s') - V^*(s')] \right| \\ &\leq \gamma \max_a \mathbb{E}_{s' \sim p(\cdot|s,a)} |V(s') - V^*(s')|. \end{aligned}$$

The first inequality is $|\max_a f(a) - \max_a g(a)| \leq \max_a |f(a) - g(a)|$; the last is Jensen. Applying the same bound to $|V(s') - V^*(s')|$ and iterating k times,

$$|V(s) - V^*(s)| \leq \gamma^k \max_{a_0, \dots, a_{k-1}} \mathbb{E}_{s_k} |V(s_k) - V^*(s_k)| \leq \gamma^k \|V - V^*\|_\infty,$$

and since $\gamma < 1$ and $\|V - V^*\|_\infty$ is finite, the right-hand side tends to 0. Hence $|V(s) - V^*(s)| = 0$ for every s . $\square$

Slide 54 illustrates the shrinking bound $\gamma^h \|V - V^*\|_\infty$ as the number of future steps h grows (its example uses $\gamma = 0.92$ and a baseline $\|V - V^*\|_\infty = 1$).




Figure 35: The bound $\gamma^h \|V - V^*\|_\infty$ for $\gamma = 0.92$, redrawn from the idea of slide 54 (values computed here).




Figure 36: Shrinking bound

Intuition

Imagine a “fake” value function V that satisfies the optimality equation and is not V^* . Each time we expand the equation one step, the gap can be at most γ times what it was; after k steps it is at most γ^k times the original gap, which melts away (with $\gamma = 0.92$, $0.92^{56} < 0.01$). The only way for the gap to survive all expansions is to be 0. Slide 55 states the consequence: *we can focus on one step at a time*, leaving the rest of the problem to $V(s')$, and any V satisfying that one-step formula is in fact V^* .

X Summary

Object / result	Formula	What it tells us	Slides
Markov state	$p(s_{t+1} s_t, a_t) = p(s_{t+1} h_t, a_t)$	the present summarises the past	15–17
MDP	$(\mathcal{S}, \mathcal{A}, p, r, \gamma)$	sequential decisions with consequences	18–29
V^π, Q^π	$\mathbb{E}[\sum_h \gamma^h r_{t+h} \mid \cdot]$	how good a state / action is under π	30–31
Finite values	$0 \leq V^\pi \leq R_{\max}/(1 - \gamma)$	why $\gamma < 1$	32
Bellman expectation	$V^\pi(s) = \mathbb{E}_a[r + \gamma \mathbb{E}_{s'} V^\pi(s')]$	one-step consistency of a given policy	33–36
$d_{s_0}^\pi$	$(1 - \gamma) \sum_h \gamma^h \mathbb{P}_h^\pi$	discounted occupancy, a true distribution	37–38
Optimal V^*, Q^*	$\sup_\pi V^\pi, \sup_\pi Q^\pi$	a deterministic policy dominates all others	39–40
Bellman optimality	$V^*(s) = \max_a[r + \gamma \mathbb{E}_{s'} V^*(s')]$	one-step lookahead with a max	41–45
Theorem 1	$\hat{\pi} \in \arg \max_a Q^* \Rightarrow V^{\hat{\pi}} = V^*$	greedy w.r.t. Q^* is optimal	46–49
Theorem 2	$V = \max_a[r + \gamma \mathbb{E}V] \Rightarrow V = V^*$	V^* is the unique solution	50–55




Figure 37: Roadmap of the lecture: from the Markov assumption to optimal policies.

Key take-aways

- An MDP is a contextual bandit whose next context depends on the action taken; the Markov property means the current state is all the agent needs to know.
- The discount factor $\gamma < 1$ keeps the value finite, with $0 \leq V^\pi \leq R_{\max}/(1 - \gamma)$, and $1/(1 - \gamma)$ is the effective horizon ($\gamma = 0.99 \Rightarrow \sim 100$ steps).
- Bellman equations turn an infinite-horizon sum into a one-step relation: value now = reward + discounted value of the next state (average for a policy, max for optimality).
- Theorem 1: acting greedily with respect to Q^* is optimal, so a deterministic stationary policy always suffices.
- Theorem 2: V^* is the *only* solution of the Bellman optimality equation, thanks to the γ -contraction; any V that satisfies the one-step formula is V^* .

A Appendix: supplementary material (in the companion PDF, not on the slides)

The following results appear in the companion PDF but have no corresponding slide; they are *not* part of the slide flow above.

A.1 Discounted sums as expectations under $d_{s_0}^\pi$

Lemma A.1 (Second part of Lemma 5 of the PDF). *For every bounded f ,*

$$\mathbb{E}^\pi \left[\sum_{h=0}^{\infty} \gamma^h f(s_h, a_h) \mid s_0 \right] = \frac{1}{1-\gamma} \mathbb{E}_{(s,a) \sim d_{s_0}^\pi} [f(s, a)].$$

Proof. It is the same exchange of (absolutely convergent) sums used in the proof of Lemma R.2, applied to f : $\mathbb{E}^\pi[\sum_h \gamma^h f(s_h, a_h) \mid s_0] = \sum_h \gamma^h \sum_{s,a} \mathbb{P}_h^\pi(s, a; s_0) f(s, a) = \frac{1}{1-\gamma} \sum_{s,a} d_{s_0}^\pi(s, a) f(s, a)$. $\square$

Intuition

Choosing $f = r$ gives $V^\pi(s_0) = \frac{1}{1-\gamma} \mathbb{E}_{(s,a) \sim d_{s_0}^\pi} [r(s, a)]$: the value is the average reward under the discounted occupancy, scaled by the effective horizon $1/(1-\gamma)$. For $\gamma = 0.9$, if the average reward under d^π is 0.5, the value is $0.5/0.1 = 5$.

A.2 Bellman operators (notation used in the PDF's proofs)

The proof of Theorem V.1 in the companion text uses two operators acting on bounded functions $V : \mathcal{S} \rightarrow \mathbb{R}$:

$$\begin{aligned} (\mathcal{B}V)(s) &= \max_a \left[r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} V(s') \right], \\ (\mathcal{B}^\pi V)(s) &= r(s, \pi(s)) + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} V(s') \quad (\pi \text{ deterministic}). \end{aligned}$$

With this notation Theorem T.1 reads $V^* = \mathcal{B}V^*$, Theorem Q.1 reads $V^\pi = \mathcal{B}^\pi V^\pi$, and Theorem W.1 says that $\mathcal{B}$ has a unique fixed point. The slides never name these operators.